

ASSIGNMENT ADVANCED HBASE 2

Case Study Description:

Let us take up the CUSTOMER and TRANSACTIONS table we have created in the Let's Do Together section. Let us solve the following use cases using these tables :-

SOLUTION:

EXPLANATION: SO HERE WE HAVE CREATED TWO TABLES "CUSTOMER" AND "TRANSACTIONS". SO FOR THAT WE NEED TO CREATE A DATABASE IN HIVE AND THEN CREATE THE TABLES. SO BELOW WE CREATED A DATABASE "acadgilddb" AND THEN USING THE SAME

COMMAND: create database acadgilddb;

SOLUTION REPORT:

```
hive> create database acadgilddb;
OK
Time taken: 0.67 seconds
```

COMMAND: use acadgilddb;

```
hive> use acadgilddb;
OK
Time taken: 0.055 seconds
```

COMMAND: CREATE TABLE CUSTOMER(custid INT, fname STRING, lname STRING, age INT, profession STRING)
row format delimited
fields terminated by ',';

EXPLANATION: CREATED A TABLE 'CUSTOMER' WITH SCHEMA FOR CUSTOMER ID, FIRST NAME, LAST NAME, AGE, PROFESSION. THEN LATER WE NEED TO LOAD THE DATASET INSIDE THE TABLE WHICH WILL HAVE ROWS DELIMITED BY '.' .

SOLUTION REPORT:

```
hive> CREATE TABLE CUSTOMER(
    > custid INT,
    > fname STRING,
    > lname STRING,
    > age INT,
    > profession STRING)
    > row format delimited fields terminated by ',';
OK
Time taken: 2.249 seconds
```

EXPLANATION: THE DATASET 'customer.txt' IS SAVED IN LOCAL DIRECTORY AT THE PATH '/home/acadgild/customer.txt' . THEN USING THE LOAD COMMAND WE WILL LOAD THE DATASET INTO THE CUSTOMER TABLE.

COMMAND:load data local inpath 'file:///home/acadgild/customer.txt' into table CUSTOMER.

SOLUTION REPORT:

```
hive> load data local inpath 'file:///home/acadgild/customer.txt' into
table CUSTOMER;
Loading data to table acadgilddb.customer
OK
Time taken: 5.711 seconds
```

EXPLANATION: CREATED A TABLE 'TRANSACTIONS' WITH SCHEMA FOR TRANSACTION NUMBER, TRANSACTION DATE, CUSTOMER ID, AMOUNT, CATEGORY, PRODUCT, CITY, STATE, SPENT BY. THEN LATER WE NEED TO LOAD THE DATASET INSIDE THE TABLE WHICH WILL HAVE ROWS DELIMITED BY '.' .

```
COMMAND: CREATE TABLE TRANSACTIONS(txnno INT, txndate STRING, custno INT,
amount DOUBLE, category STRING, product STRING,          city STRING,
state STRING, spendby STRING)
        row format delimited
        fields terminated by ',';
```

SOLUTION:

```
hive> CREATE TABLE TRANSACTIONS (
    > txnno INT,
    > txndate STRING,
    > custno INT,
    > amount DOUBLE,
    > category STRING,
    > product STRING,
    > city STRING,
    > state STRING,
    > spendby STRING)
    > row format delimited fields terminated by ',';
```

OK

Time taken: 0.358 seconds

EXPLANATION: THE DATASET 'transaction.txt' IS SAVED IN LOCAL DIRECTORY AT THE PATH '/home/acadgild/transaction.txt' . THEN USING THE LOAD COMMAND WE WILL LOAD THE DATASET INTO THE CUSTOMER TABLE.

```
COMMAND:load data local inpath 'file:///home/acadgild/transaction.txt'
into table TRANSACTIONS;
```

SOLUTION REPORT:

```
hive> load data local inpath 'file:///home/acadgild/transaction.txt' into
table TRANSACTIONS;
Loading data to table acadgilddb.transactions
OK
Time taken: 2.162 seconds
```

EXPLANATION: USING THE SELECT OPERATION WE CAN SEE THE DATASET HAS BEEN LOADED TO THE CUSTOMER TABLE AND TRANSACTIONS TABLE.

SOLUTION REPORT:

```
hive> SELECT * FROM CUSTOMER;
```

```

OK
101  Amitabh      Bacchan      65      Actor
102  Sharukh      Khan      45      Doctor
103  Akshay      Kumar      38      Dentist
104  Anubahv      kumar      58      Business
105  Pawan Trivedi      34      service
106  Aamir Null      42      scientest
107  Salman      Khan      43      Surgen
108  Ranbir      Kapoor      26      Industrialist
NULL NULL NULL NULL NULL
Time taken: 0.394 seconds, Fetched: 9 row(s)

```

```

hive> SELECT * FROM TRANSACTIONS;
OK
97834 05/02/2018 101 965.0 Entertainment Movie Pune Maharashtra
      Daughter
98396 12/01/2018 102 239.0 Food Grocery Patna Bihar Self
34908 06/01/2018 101 875.0 Travel Air Bangalore Karnataka
      Spouse
70958 17/02/2018 104 439.0 Food Restaurant Delhi Delhi Wife
9874 21/01/2018 105 509.0 Entertainment Park Kolkata West
Bengal NULL
94585 19/01/2018 106 629.0 Rent House Hyderabad Telangana Self
45509 20/01/2018 107 953.0 Travel Rail Chennai Tamil Nadu
      Brother
7864 01/02/2018 108 569.0 Rent Parking Goa Goa Wife
NULL NULL NULL NULL NULL NULL NULL NULL NULL
Time taken: 0.66 seconds, Fetched: 9 row(s)

```

TASK 1: Find out the number of transaction done by each customer
(These should be take up in module 8 itself)

COMMAND: SELECT CUSTID, FNAME, SUM(AMOUNT) FROM CUSTOMER C, TRANSACTIONS
T WHERE C.CUSTID=T.CUSTNO GROUP BY CUSTID, FNAME;

EXPLANATION: HERE THIS OPERATION GIVES THE RESULT WITH THE NUMBER OF
TRANSACTIONS DONE BY EACH CUSTOMER. SO THE SELECT OPERATION SELECTS THE
CUSTOMER ID, FIRST AND THE TOTAL OF THE AMOUNT FROM THE TABLES CUSTOMER
AND TRANSACTIONS WHICH ARE JOINED TOGETHER USING 'JOIN' KEYWORD WITH
ALIAS NAME 'C' FOR CUSTOMER AND 'T' FOR TRANSACTIONS.

SOLUTION REPORT:

```

hive> SELECT CUSTID, FNAME, SUM(AMOUNT) FROM CUSTOMER C, TRANSACTIONS T
WHERE C.CUSTID=T.CUSTNO GROUP BY CUSTID, FNAME;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in
the future versions. Consider using a different execution engine (i.e.
spark, tez) or using Hive 1.X releases.
Query ID = acadgild_20190109035131_6c758d49-1a24-4aac-95eb-52d919d64846
Total jobs = 1

```

```

SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/acadgild/install/hive/apache-
hive-2.3.2-bin/lib/log4j-slf4j-impl-
2.6.2.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/acadgild/install/hadoop/hadoop-
2.6.5/share/hadoop/common/lib/slf4j-log4j12-
1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an
explanation.
SLF4J: Actual binding is of type
[org.apache.logging.slf4j.Log4jLoggerFactory]
2019-01-09 03:51:52      Starting to launch local task to process map join;
      maximum memory = 518979584
2019-01-09 03:51:56      Dump the side-table for tag: 0 with group count: 8
into file: file:/tmp/acadgild/35c7bbfa-9569-48e6-9072-
77162e214a97/hive_2019-01-09_03-51-31_385_5914681181293595335-1/-local-
10005/HashTable-Stage-2/MapJoin-mapfile120--.hashtable
2019-01-09 03:51:56      Uploaded 1 File to: file:/tmp/acadgild/35c7bbfa-
9569-48e6-9072-77162e214a97/hive_2019-01-09_03-51-
31_385_5914681181293595335-1/-local-10005/HashTable-Stage-2/MapJoin-
mapfile120--.hashtable (469 bytes)
2019-01-09 03:51:56      End of local task; Time Taken: 4.835 sec.
Execution completed successfully
MapredLocal task succeeded
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
    set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
    set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
    set mapreduce.job.reduces=<number>
Starting Job = job_1546964532394_0014, Tracking URL =
http://localhost:8088/proxy/application_1546964532394_0014/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job
-kill job_1546964532394_0014
Hadoop job information for Stage-2: number of mappers: 1; number of
reducers: 1
2019-01-09 03:52:19,191 Stage-2 map = 0%,   reduce = 0%
2019-01-09 03:52:39,964 Stage-2 map = 100%,   reduce = 0%, Cumulative CPU
6.29 sec
2019-01-09 03:53:01,987 Stage-2 map = 100%,   reduce = 67%, Cumulative CPU
10.62 sec
2019-01-09 03:53:04,568 Stage-2 map = 100%,   reduce = 100%, Cumulative
CPU 11.83 sec
MapReduce Total cumulative CPU time: 11 seconds 830 msec
Ended Job = job_1546964532394_0014
MapReduce Jobs Launched:
Stage-Stage-2: Map: 1   Reduce: 1   Cumulative CPU: 11.83 sec   HDFS Read:
13364 HDFS Write: 292 SUCCESS
Total MapReduce CPU Time Spent: 11 seconds 830 msec
OK
101   Amitabh       1840.0
102   Sharukh       239.0

```

```
104  Anubahv      439.0
105  Pawan 509.0
106  Aamir 629.0
107  Salman      953.0
108  Ranbir      569.0
Time taken: 95.689 seconds, Fetched: 7 row(s)
```

TASK 2: Create a new table called TRANSACTIONS_COUNT. This table should have 3 fields - custid, fname and count. (Again to be done in module 8)

```
COMMAND: create table TRANSACTION_COUNT(custid INT, fname STRING, count
INT)
        row format delimited
        fields terminated by ',';
```

EXPLANATION: CREATED TABLE TRANSACTIONS_COUNT AND PROVIDED SCHEMA FOR THE FIELDS CUSTOMER ID, FIRST NAME AND COUNT.

SOLUTION REPORT:

```
hive> create table TRANSACTION_COUNT(custid INT, fname STRING, count INT)
      > row format delimited
      > fields terminated by ',';
OK
Time taken: 0.219 seconds
```

TASK 3: Now write a hive query in such a way that the query populates the data obtained in Step 1 above and populate the table in step 2 above. (This has to be done in module 9).

```
COMMAND: INSERT INTO TABLE TRANSACTION_COUNT SELECT CUSTID, FNAME,
COUNT(DISTINCT CUSTID) FROM CUSTOMER A JOIN TRANSACTIONS B WHERE
A.CUSTID=B.CUSTNO GROUP BY CUSTID,FNAME;
```

EXPLANATION: USING THE INSERT COMMAND WE ARE LOADING THE TABLE WITH DATA PRODUCED IN THE OUTPUT OF TASK 1. SO HERE WE USING A SELECT OPERATION TO SAVE CUSTOMER ID, FIRST NAME AND COUNT FROM THE COMBINED TABLE(CUSTOMER AND TRANSACTIONS) IN KEYWORD "COUNT(DISTICT CUSTID)" WILL GIVE THE NUMBER OF CUSTOMER ID'S AND THEN INSERT KEYWORD WILL ENTER THE FOLLOWING DATA INTO THE TABLE UNDER FIELDS CUSTID, FNAME AND COUNT.

SOLUTION REPORT:

```
hive> INSERT INTO TABLE TRANSACTION_COUNT SELECT CUSTID, FNAME,
COUNT(DISTINCT CUSTID) FROM CUSTOMER A JOIN TRANSACTIONS B WHERE
A.CUSTID=B.CUSTNO GROUP BY CUSTID,FNAME;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in
the future versions. Consider using a different execution engine (i.e.
spark, tez) or using Hive 1.X releases.
Query ID = acadgild_20190109021312_acc018a2-685e-4d27-99e6-ea66495b5ed0
Total jobs = 1
```

SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/acadgild/install/hive/apache-hive-2.3.2-bin/lib/log4j-slf4j-impl-2.6.2.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
2019-01-09 02:13:23 Starting to launch local task to process map join;
maximum memory = 518979584
2019-01-09 02:13:26 Dump the side-table for tag: 0 with group count: 8
into file: file:/tmp/acadgild/35c7bbfa-9569-48e6-9072-77162e214a97/hive_2019-01-09_02-13-12_128_8238481662940432256-1/-local-10003/HashTable-Stage-2/MapJoin-mapfile90--.hashtable
2019-01-09 02:13:26 Uploaded 1 File to: file:/tmp/acadgild/35c7bbfa-9569-48e6-9072-77162e214a97/hive_2019-01-09_02-13-12_128_8238481662940432256-1/-local-10003/HashTable-Stage-2/MapJoin-mapfile90--.hashtable (469 bytes)
2019-01-09 02:13:26 End of local task; Time Taken: 2.88 sec.
Execution completed successfully
MapredLocal task succeeded
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
set mapreduce.job.reduces=<number>
Starting Job = job_1546964532394_0011, Tracking URL =
http://localhost:8088/proxy/application_1546964532394_0011/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1546964532394_0011
Hadoop job information for Stage-2: number of mappers: 1; number of reducers: 1
2019-01-09 02:13:40,210 Stage-2 map = 0%, reduce = 0%
2019-01-09 02:13:51,337 Stage-2 map = 100%, reduce = 0%, Cumulative CPU 2.95 sec
2019-01-09 02:14:05,726 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 6.69 sec
MapReduce Total cumulative CPU time: 6 seconds 690 msec
Ended Job = job_1546964532394_0011
Loading data to table acadgilddb.transaction_count
MapReduce Jobs Launched:
Stage-Stage-2: Map: 1 Reduce: 1 Cumulative CPU: 6.69 sec HDFS Read: 13962 HDFS Write: 176 SUCCESS
Total MapReduce CPU Time Spent: 6 seconds 690 msec
OK
Time taken: 55.538 seconds

EXPLANATION: USING THE SELECT OPERATION WE CAN SEE THE TRANSACTION_COUNT HAS BEEN LOADED WITH THE REQUIRED DATA.

SOLUTION REPORT:

```
hive> SELECT * FROM TRANSACTION_COUNT;
OK
101  Amitabh      1
102  Sharukh      1
104  Anubahv      1
105  Pawan 1
106  Aamir 1
107  Salman       1
108  Ranbir       1
Time taken: 0.201 seconds, Fetched: 7 row(s)
```

TASK 4: Now let's make the TRANSACTIONS_COUNT table Hbase compliant. In the sense, use Ser Des And Storage handler features of hive to change the TRANSACTIONS_COUNT table to be able to create a TRANSACTIONS table in Hbase. (This has to be done in module 10)

EXPLANATION: SO FOR DOING THE ABOVE OPERATION, FIRST LOGIN INTO HBASE SHELL AND THEN CREATE A TRANSACTIONS TABLE IN HBASE WITH A COLUMN FAMILY "DETAILS".

COMMAND: create 'TRANSACTIONS','DETAILS'

SOLUTION REPORT:

```
hbase(main):003:0> create 'TRANSACTIONS','DETAILS'
0 row(s) in 2.8030 seconds
```

=> Hbase::Table - TRANSACTIONS

EXPLANATION: NOW WE CREATED A HIVE TABLE POINTING TO HBASE TABLE. HERE CREATED A EXTERNAL TABLE "TRANSACTIONS_hbase" WITH THE SCHEMA FOR THE FIELDS CUSTOMER ID, FIRST NAME AND COUNT. AS WE CREATED A NON-NATIVE HIVE TABLE USING STORAGE HANDLER SO USED THE STORED BY clause. AS FOR THE INTEGRATING HBASE WITH HIVE STORAGE HANDLERS IN HIVE ARE USED. THEY ARE THE COMBINATION OF INPUT AND OUTPUT FORMAT AS SERDE A SPECIFIC CODE THAT HIVE USES TO IDENTIFY AN EXTERNAL ENTITY AS A HIVE TABLE. HERE "hbase.columns.mapping" IS USED TO MAP THE HIVE COLUMNS WITH THE HBASE COLUMNS. THE FIRST COLUMN MUST BE THE KEY COLUMN WHICH WOULD ALSO BE SAME AS HBASE'S ROW KEY COLUMN AND PROVIDED THE HBASE TABLE "TRANSACTIONS" SO THAT THE HIVE TABLE CONNECTS WITH THE HBASE TABLE.

```
COMMAND: create external table TRANSACTIONS_hbase(custid INT, fname
string, count INT)
        STORED BY 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'
        with serdeproperties
("hbase.columns.mapping"=":key,DETAILS:fname,DETAILS:count")
        tblproperties("hbase.table.name"="TRANSACTIONS");
```

SOLUTION REPORT:

```
hive> create external table TRANSACTIONS_hbase(custid INT, fname string,
count INT)
> STORED BY 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'
> with serdeproperties
("hbase.columns.mapping"=":key,DETAILS:fname,DETAILS:count")
> tblproperties("hbase.table.name"="TRANSACTIONS");
OK
Time taken: 1.947 seconds
```

TASK 5: Now insert the data in TRANSACTIONS_COUNT table using the query in step 3 again, this should populate the Hbase TRANSACTIONS table automatically (This has to be done in module 10)

EXPLANATION: HERE USED THE SAME QUERY WHICH WAS USED IN TASK 3, TO INSERT THE DATA INTO THE TABLE "TRANSACTIONS_hbase" WHICH WAS CREATED ABOVE AS A EXTERNAL TABLE SO THAT THE DATA IS INSERTED INTO HIVE TABLE AND LATER WHEN THE QUERY IS EXECUTED THE DATA IS AUTOMATICALLY LOADED INTO HBASE TABLE "TRANSACTIONS" AS WELL.

COMMAND: INSERT INTO TABLE TRANSACTIONS_hbase SELECT CUSTID, FNAME, COUNT(DISTINCT CUSTID) FROM CUSTOMER A JOIN TRANSACTIONS B WHERE A.CUSTID=B.CUSTNO GROUP BY CUSTID,FNAME;

SOLUTION REPORT:

```
hive> INSERT INTO TABLE TRANSACTIONS_hbase SELECT CUSTID, FNAME,
COUNT(DISTINCT CUSTID) FROM CUSTOMER A JOIN TRANSACTIONS B WHERE
A.CUSTID=B.CUSTNO GROUP BY CUSTID,FNAME;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in
the future versions. Consider using a different execution engine (i.e.
spark, tez) or using Hive 1.X releases.
Query ID = acadgild_20190109025722_20bbbd3a-64cc-4704-867d-bc3aa7fe458f
Total jobs = 1
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/acadgild/install/hive/apache-
hive-2.3.2-bin/lib/log4j-slf4j-impl-
2.6.2.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/acadgild/install/hadoop/hadoop-
2.6.5/share/hadoop/common/lib/slf4j-log4j12-
1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an
explanation.
SLF4J: Actual binding is of type
[org.apache.logging.slf4j.Log4jLoggerFactory]
2019-01-09 02:57:43 Starting to launch local task to process map join;
maximum memory = 518979584
2019-01-09 02:57:49 Dump the side-table for tag: 0 with group count: 8
into file: file:/tmp/acadgild/35c7bbfa-9569-48e6-9072-
77162e214a97/hive_2019-01-09_02-57-22_059_2728555604714371678-1/-local-
10002/HashTable-Stage-4/MapJoin-mapfile110--.hashtable
2019-01-09 02:57:49 Uploaded 1 File to: file:/tmp/acadgild/35c7bbfa-
9569-48e6-9072-77162e214a97/hive_2019-01-09_02-57-
22_059_2728555604714371678-1/-local-10002/HashTable-Stage-4/MapJoin-
mapfile110--.hashtable (469 bytes)
2019-01-09 02:57:49 End of local task; Time Taken: 5.698 sec.
```


Execution completed successfully
MapredLocal task succeeded
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
 set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
 set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
 set mapreduce.job.reduces=<number>
Starting Job = job_1546964532394_0013, Tracking URL =
http://localhost:8088/proxy/application_1546964532394_0013/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job
-kill job_1546964532394_0013
Hadoop job information for Stage-4: number of mappers: 1; number of
reducers: 1
2019-01-09 02:58:22,299 Stage-4 map = 0%, reduce = 0%
2019-01-09 02:58:42,899 Stage-4 map = 100%, reduce = 0%, Cumulative CPU
5.99 sec
2019-01-09 02:59:06,766 Stage-4 map = 100%, reduce = 67%, Cumulative CPU
10.62 sec
2019-01-09 02:59:11,512 Stage-4 map = 100%, reduce = 100%, Cumulative
CPU 14.47 sec
MapReduce Total cumulative CPU time: 14 seconds 470 msec
Ended Job = job_1546964532394_0013
MapReduce Jobs Launched:
Stage-Stage-4: Map: 1 Reduce: 1 Cumulative CPU: 14.47 sec HDFS Read:
20550 HDFS Write: 0 SUCCESS
Total MapReduce CPU Time Spent: 14 seconds 470 msec
OK
Time taken: 111.756 seconds

EXPLANATION: HERE USING SELECT OPERATION WE CAN SEE THE DATA HAS BEEN
INSERTED INTO THE HIVE TABLE "TRANSACTIONS_hbase".

SOLUTION REPORT:

```
hive> select * from TRANSACTIONS_hbase;
OK
101  Amitabh      1
102  Sharukh      1
104  Anubahv      1
105  Pawan        1
106  Aamir        1
107  Salman       1
108  Ranbir       1
Time taken: 0.921 seconds, Fetched: 7 row(s)
```

TASK 6: Now from the Hbase level, write the Hbase java API code to access
and scan the TRANSACTIONS table data from java level.

COMMAND: scan 'TRANSACTIONS'

EXPLANATION: TO CHECK WHETHER THE DATA FROM THE HIVE TABLE
"TRANSACTIONS_hbase" HAS BEEN LOADED INTO THE HBASE TABLE
"TRANSACTIONS".SO USED THE BELOW COMMAND.

SOLUTION REPORT:

hbase(main):004:0> scan 'TRANSACTIONS'

ROW	COLUMN+CELL
101	column=DETAILS:count, timestamp=1546982950134,
value=1	
101	column=DETAILS:fname, timestamp=1546982950134,
value=Amita	
	bh
102	column=DETAILS:count, timestamp=1546982950134,
value=1	
102	column=DETAILS:fname, timestamp=1546982950134,
value=Sharu	
	kh
104	column=DETAILS:count, timestamp=1546982950134,
value=1	
104	column=DETAILS:fname, timestamp=1546982950134,
value=Anuba	
	hv
105	column=DETAILS:count, timestamp=1546982950134,
value=1	
105	column=DETAILS:fname, timestamp=1546982950134,
value=Pawan	
106	column=DETAILS:count, timestamp=1546982950134,
value=1	
106	column=DETAILS:fname, timestamp=1546982950134,
value=Aamir	
107	column=DETAILS:count, timestamp=1546982950134,
value=1	
107	column=DETAILS:fname, timestamp=1546982950134,
value=Salma	
	n
108	column=DETAILS:count, timestamp=1546982950134,
value=1	
108	column=DETAILS:fname, timestamp=1546982950134,
value=Ranbi	

7 row(s) in 2.0580 seconds

OUTPUT

TASK 1:

```
101    Amitabh 1840.0
102    Sharukh 239.0
104    Anubahv 439.0
105    Pawan   509.0
106    Aamir   629.0
107    Salman  953.0
108    Ranbir  569.0
Time taken: 95.689 seconds, Fetched: 7 row(s)
```

TASK 3:

```
hive> SELECT * FROM TRANSACTION_COUNT;
OK
101    Amitabh 1
102    Sharukh 1
104    Anubahv 1
105    Pawan   1
106    Aamir   1
107    Salman  1
108    Ranbir  1
Time taken: 0.201 seconds, Fetched: 7 row(s)
```

TASK 5:

```
hive> select * from TRANSACTIONS_hbase;
OK
101    Amitabh 1
102    Sharukh 1
104    Anubahv 1
105    Pawan   1
106    Aamir   1
107    Salman  1
108    Ranbir  1
Time taken: 3.151 seconds, Fetched: 7 row(s)
```

TASK 6:

```
hbase(main):004:0> scan 'TRANSACTIONS'
```

ROW	COLUMN+CELL
101	column=DETAILS:count, timestamp=1546982950134, value=1
101	column=DETAILS:fname, timestamp=1546982950134, value=Amita bh
102	column=DETAILS:count, timestamp=1546982950134, value=1
102	column=DETAILS:fname, timestamp=1546982950134, value=Sharu kh
104	column=DETAILS:count, timestamp=1546982950134, value=1
104	column=DETAILS:fname, timestamp=1546982950134, value=Anuba hv
105	column=DETAILS:count, timestamp=1546982950134, value=1
105	column=DETAILS:fname, timestamp=1546982950134, value=Pawan
106	column=DETAILS:count, timestamp=1546982950134, value=1
106	column=DETAILS:fname, timestamp=1546982950134, value=Aamir
107	column=DETAILS:count, timestamp=1546982950134, value=1
107	column=DETAILS:fname, timestamp=1546982950134, value=Salma n
108	column=DETAILS:count, timestamp=1546982950134, value=1
108	column=DETAILS:fname, timestamp=1546982950134, value=Ranbi r

```
7 row(s) in 2.0580 seconds
```